



# COMP4211 Final Project

## Project 1

Jasper | Ryder | Yerke | Daniel






# The Problem

What needs to be done?

# Design Goals


---




1. Robust
2. Good Precision
3. Good Recall
4. Good F1 Score

# Data Augmentation

Original Data



Augmented Data



# The First Task - Overview

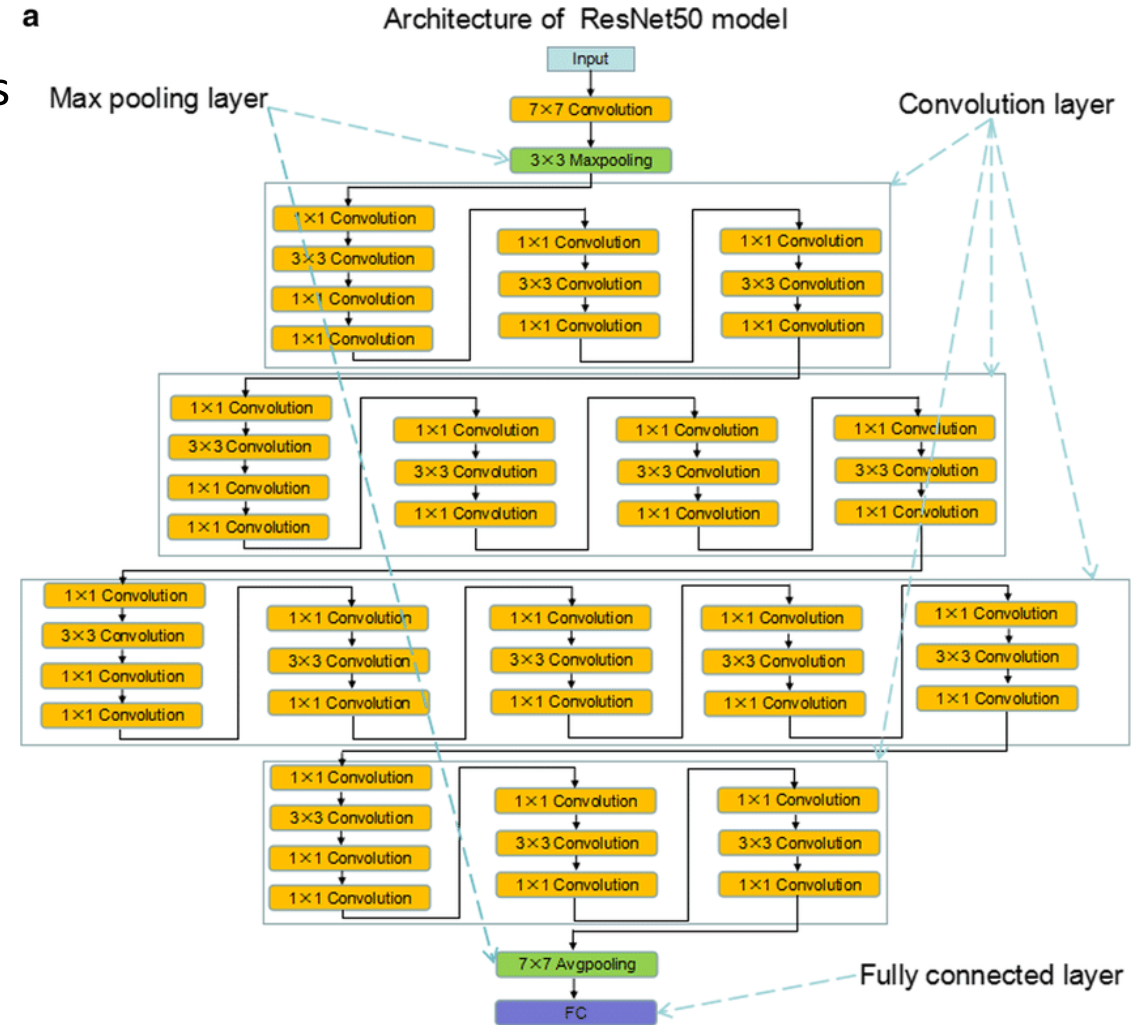
- To build a CNN model using cross-entropy loss

## Model Specifications

- ResNet50 Model
- Stochastic Gradient Descent:
  - Learning Rate = 0.001
  - Momentum = 0.9
- Uses cross-entropy loss function

$$\begin{aligned} L(x)_{CE} &= - \sum_{j=1}^n t_j \log(p_j) \\ &= -t_i \log(p_i) \\ &= -\log(p_i) \end{aligned}$$

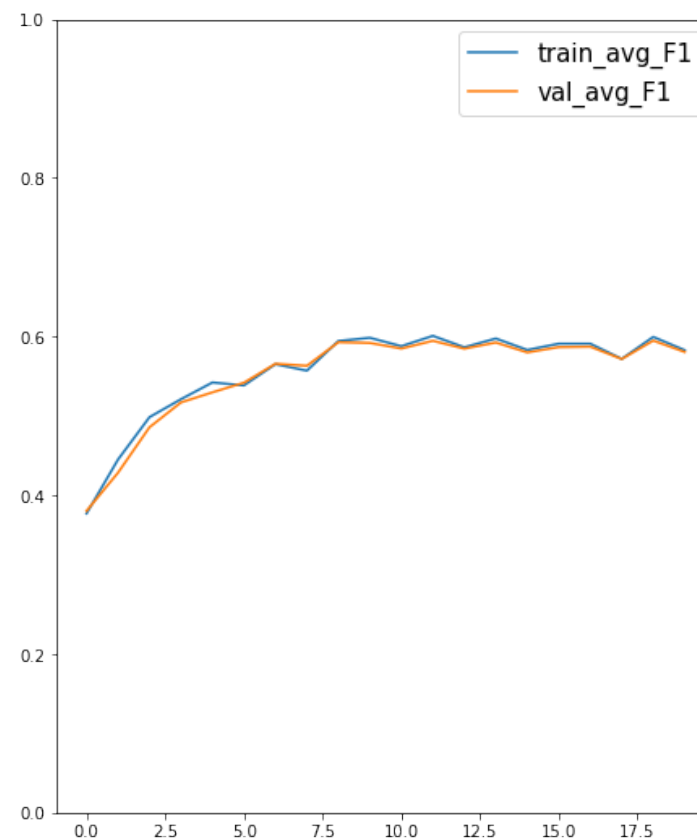

Where  $t_i :=$  the basis vector  $e_i \in \mathbb{R}^n$  if the label of  $x$  is  $i$ .



Source: [https://www.researchgate.net/figure/The-architecture-of-ResNet50-and-deep-learning-model-flowchart-a-b-Architecture-of\\_fig1\\_334767096](https://www.researchgate.net/figure/The-architecture-of-ResNet50-and-deep-learning-model-flowchart-a-b-Architecture-of_fig1_334767096)

## The First Task - Results

	Precision (%)	Recall (%)	F1 (%)
Class 1 (NIL)	86.14	76.37	80.29
Class 2 (Moderate)	64.51	92.76	75.54
Class 3 (Severe)	66.63	14.63	22.64



# The Second Task - Overview

## Model Specifications

- ResNet50 Model (same as Task 1)
- Uses LDAM loss function

$$\mathcal{L}_{\text{LDAM}}((x, y); f) = -\log \frac{e^{z_y - \Delta_y}}{e^{z_y - \Delta_y} + \sum_{j \neq y} e^{z_j}}$$

where  $\Delta_j = \frac{C}{n_j^{1/4}}$  for  $j \in \{1, \dots, k\}$

```
def ldam(out, label, C, n_j, drw_active):
    deltas = 1/n_j # actually = 1/(n_j)^0.25

    # for a given input x with corresponding label y and model-generated output
    # which is a vector (p_0, p_1, p_2) where p_i denotes the probability the
    # model believes x belongs to class i
    # we need to extract the necessary values into a B*1 vector z_y. Example:
    # (0.4, 0.32, 0.28) | 1
    # (0.8, 0.15, 0.05) | 2
    # (0.2, 0.6, 0.2) | 1

    z_y = out[range(out.shape[0]), label]
    # returns vector of size B*1
    # corresponding to z_y of each element x in the current batch

    delta_y = torch.zeros((label.shape[0],)).cuda() # let delta_y be B*1
    for i in range(3): # since k=3. Flexible if k changes
        delta_y += torch.where(label+1==i+1, label+1, 0)*deltas[i]/(i+1)
    delta_y_copy = delta_y.clone() # used for DRW
    delta_y *= C
    e_zyminusdeltay = torch.exp(z_y-delta_y)

    # since each value of label is in {0, 1, 2}
    z_j_1 = out[range(out.shape[0]), (label+1)%3]
    z_j_2 = out[range(out.shape[0]), (label+2)%3]
    e_zj_sum = torch.exp(z_j_1) + torch.exp(z_j_2)

    # compute L_LDAM((x,y);f)
    # for all x in the batch. Asked TA and he said it was nat log
    result = -torch.log(e_zyminusdeltay/(e_zyminusdeltay+e_zj_sum))

    # reweight LDAM via DRW if necessary via entrywise tensor multiplication
    renormalize_factor = 0
    if drw_active:
        n_y_reciprocal = torch.pow(delta_y_copy, 4)
        result = result * n_y_reciprocal
        renormalize_factor = torch.sum(n_y_reciprocal).item()

    # take average of the entries to get a real-valued result
    return torch.mean(result), renormalize_factor
```



## The Second Task - Overview

$n_j$ : number of samples in training set who have label  $j$ , if 591 samples have label 0, 839 have label 1, 324 have label 2

$\text{deltas} = [1/(591^{0.25}), 1/(839^{0.25}), 1/(324^{0.25})]$

```
def ldam(out, label, C, n_j, drw_active):
    deltas = 1/n_j # actually = 1/(n_j)^0.25

    # for a given input x with corresponding label y and model-generated output
    # which is a vector (p_0, p_1, p_2) where p_i denotes the probability the
    # model believes x belongs to class i
    # we need to extract the necessary values into a B*1 vector z_y. Example:
    # (0.4, 0.32, 0.28) | 1
    # (0.8, 0.15, 0.05) | 2
    # (0.2, 0.6, 0.2) | 1

    z_y = out[range(out.shape[0]), label]
    # returns vector of size B*1
    # corresponding to z_y of each element x in the current batch

    delta_y = torch.zeros((label.shape[0],)).cuda() # let delta_y be B*1
    for i in range(3): # since k=3. Flexible if k changes
        delta_y += torch.where(label+1==i+1, label+1, 0)*deltas[i]/(i+1)
    delta_y_copy = delta_y.clone() # used for DRW
    delta_y *= C
    e_zyminusdeltay = torch.exp(z_y-delta_y)

    # since each value of label is in {0, 1, 2}
    z_j_1 = out[range(out.shape[0]), (label+1)%3]
    z_j_2 = out[range(out.shape[0]), (label+2)%3]
    e_zj_sum = torch.exp(z_j_1) + torch.exp(z_j_2)

    # compute L_LDAM((x,y);f)
    # for all x in the batch. Asked TA and he said it was nat log
    result = -torch.log(e_zyminusdeltay/(e_zyminusdeltay+e_zj_sum))

    # reweight LDAM via DRW if necessary via entrywise tensor multiplication
    renormalize_factor = 0
    if drw_active:
        n_y_reciprocal = torch.pow(delta_y_copy, 4)
        result = result * n_y_reciprocal
        renormalize_factor = torch.sum(n_y_reciprocal).item()

    # take average of the entries to get a real-valued result
    return torch.mean(result), renormalize_factor
```



## The Second Task - Overview

$n_j$ : number of samples in training set who have label  $j$ , if 591 samples have label 0, 839 have label 1, 324 have label 2

$\text{deltas} = [1/(591^{0.25}), 1/(839^{0.25}), 1/(324^{0.25})]$

$$\mathcal{L}_{\text{LDAM}}((x, y); f) = -\log \frac{e^{z_y - \Delta_y}}{e^{z_y - \Delta_y} + \sum_{j \neq y} e^{z_j}}$$

where  $\Delta_j = \frac{C}{n_j^{1/4}}$  for  $j \in \{1, \dots, k\}$

Consider some input  $x$  whose model output is (0.4, 0.32, 0.28)

Let  $y = 1$  be  $x$ 's corresponding label.

So  $z_y$  just finds (0.4, 0.32, 0.28)[1] = 0.32

Do this for every input  $x$ ...

$\text{delta}_y$  just finds  $C/(n_y)^{0.25}$  for every  $x$

```
def ldam(out, label, C, n_j, drw_active):
    deltas = 1/n_j # actually = 1/(n_j)^0.25

    # for a given input x with corresponding label y and model-generated output
    # which is a vector (p_0, p_1, p_2) where p_i denotes the probability the
    # model believes x belongs to class i
    # we need to extract the necessary values into a B*1 vector z_y. Example:
    # (0.4, 0.32, 0.28) | 1
    # (0.8, 0.15, 0.05) | 2
    # (0.2, 0.6, 0.2) | 1

    z_y = out[range(out.shape[0]), label]
    # returns vector of size B*1
    # corresponding to z_y of each element x in the current batch

    delta_y = torch.zeros((label.shape[0],)).cuda() # let delta_y be B*1
    for i in range(3): # since k=3. Flexible if k changes
        delta_y += torch.where(label+1==i+1, label+1, 0)*deltas[i]/(i+1)
    delta_y_copy = delta_y.clone() # used for DRW
    delta_y *= C
    e_zyminusdeltay = torch.exp(z_y-delta_y)

    # since each value of label is in {0, 1, 2}
    z_j_1 = out[range(out.shape[0]), (label+1)%3]
    z_j_2 = out[range(out.shape[0]), (label+2)%3]
    e_zj_sum = torch.exp(z_j_1) + torch.exp(z_j_2)

    # compute L_LDAM((x,y);f)
    # for all x in the batch. Asked TA and he said it was nat log
    result = -torch.log(e_zyminusdeltay/(e_zyminusdeltay+e_zj_sum))

    # reweight LDAM via DRW if necessary via entrywise tensor multiplication
    renormalize_factor = 0
    if drw_active:
        n_y_reciprocal = torch.pow(delta_y_copy, 4)
        result = result * n_y_reciprocal
        renormalize_factor = torch.sum(n_y_reciprocal).item()

    # take average of the entries to get a real-valued result
    return torch.mean(result), renormalize_factor
```

## The Second Task - Overview

$$\mathcal{L}_{\text{LDAM}}((x, y); f) = -\log \frac{e^{z_y - \Delta_y}}{e^{z_y - \Delta_y} + \sum_{j \neq y} e^{z_j}}$$

where  $\Delta_j = \frac{C}{n_j^{1/4}}$  for  $j \in \{1, \dots, k\}$

Consider some input  $x$  whose model output is  $(0.4, 0.32, 0.28)$   
Let  $y = 1$  be  $x$ 's corresponding label.

**e\_zj\_sum** just finds  $e^{z_j}$  for every incorrect label for  $x$  and adds them

Since only 3 labels:  $\{0, 1, 2\}$

Then the incorrect labels can always be found by  $(y+1) \bmod 3$  and  $(y+2) \bmod 3$

Do this for every input  $x$ ...

Then just plug in yellow and light blue terms into LDAM function

```
def ldam(out, label, C, n_j, drw_active):
    deltas = 1/n_j # actually = 1/(n_j)^0.25

    # for a given input x with corresponding label y and model-generated output
    # which is a vector (p_0, p_1, p_2) where p_i denotes the probability the
    # model believes x belongs to class i
    # we need to extract the necessary values into a B*1 vector z_y. Example:
    # (0.4, 0.32, 0.28) | 1
    # (0.8, 0.15, 0.05) | 2
    # (0.2, 0.6, 0.2) | 1

    z_y = out[range(out.shape[0]), label]
    # returns vector of size B*1
    # corresponding to z_y of each element x in the current batch

    delta_y = torch.zeros((label.shape[0],)).cuda() # let delta_y be B*1
    for i in range(3): # since k=3. Flexible if k changes
        delta_y += torch.where(label+1==i+1, label+1, 0)*deltas[i]/(i+1)
    delta_y_copy = delta_y.clone() # used for DRW
    delta_y *= C
    e_zyminusdeltay = torch.exp(z_y-delta_y)

    # since each value of label is in {0, 1, 2}
    z_j_1 = out[range(out.shape[0]), (label+1)%3]
    z_j_2 = out[range(out.shape[0]), (label+2)%3]
    e_zj_sum = torch.exp(z_j_1) + torch.exp(z_j_2)

    # compute L_LDAM((x,y);f)
    # for all x in the batch. Asked TA and he said it was nat log
    result = -torch.log(e_zyminusdeltay/(e_zyminusdeltay+e_zj_sum))

    # reweight LDAM via DRW if necessary via entrywise tensor multiplication
    renormalize_factor = 0
    if drw_active:
        n_y_reciprocal = torch.pow(delta_y_copy, 4)
        result = result * n_y_reciprocal
        renormalize_factor = torch.sum(n_y_reciprocal).item()

    # take average of the entries to get a real-valued result
    return torch.mean(result), renormalize_factor
```

## The Second Task - Overview

- **Hyperparameters to tweak:**
  - C
  - Learning rate
  - Momentum term
  - alpha (regularization)
  - Batch size
- Calculating the optimal set of hyperparameters is infeasible
- **Solution: exact grid search**
  - Try every possible combination of parameters
  - Too many combinations...
- Instead, opt to fix some values: use industry standard for batch size (64) momentum (0.9) and learning rate (0.001)
- alpha = (0.001, **0.01**, 0.1, 1), C = (-2, **-0.25**, 0.25, 2)

```
def ldam(out, label, C, n_j, drw_active):
    deltas = 1/n_j # actually = 1/(n_j)^0.25

    # for a given input x with corresponding label y and model-generated output
    # which is a vector (p_0, p_1, p_2) where p_i denotes the probability the
    # model believes x belongs to class i
    # we need to extract the necessary values into a B*1 vector z_y. Example:
    # (0.4, 0.32, 0.28) | 1
    # (0.8, 0.15, 0.05) | 2
    # (0.2, 0.6, 0.2) | 1

    z_y = out[range(out.shape[0]), label]
    # returns vector of size B*1
    # corresponding to z_y of each element x in the current batch

    delta_y = torch.zeros((label.shape[0],)).cuda() # let delta_y be B*1
    for i in range(3): # since k=3. Flexible if k changes
        delta_y += torch.where(label+1==i+1, label+1, 0)*deltas[i]/(i+1)
    delta_y_copy = delta_y.clone() # used for DRW
    delta_y *= C
    e_zyminusdeltay = torch.exp(z_y-delta_y)

    # since each value of label is in {0, 1, 2}
    z_j_1 = out[range(out.shape[0]), (label+1)%3]
    z_j_2 = out[range(out.shape[0]), (label+2)%3]
    e_zj_sum = torch.exp(z_j_1) + torch.exp(z_j_2)



    # compute L_LDAM((x,y);f)
    # for all x in the batch. Asked TA and he said it was nat log
    result = -torch.log(e_zyminusdeltay/(e_zyminusdeltay+e_zj_sum))

    # reweight LDAM via DRW if necessary via entrywise tensor multiplication
    renormalize_factor = 0
    if drw_active:
        n_y_reciprocal = torch.pow(delta_y_copy, 4)
        result = result * n_y_reciprocal
        renormalize_factor = torch.sum(n_y_reciprocal).item()

    # take average of the entries to get a real-valued result
    return torch.mean(result), renormalize_factor
```

## The Second Task - Results

	Precision (%)	Recall (%)	F1 (%)
Class 1 (NIL)	87.51	67.84	75.45
Class 2 (Moderate)	60.19	94.66	72.90
Class 3 (Severe)	52.18	7.56	12.80



# The Final Model

---




## Issues With the Other Models

- Overfit to majority classes in the training set
- Doesn't perform well on the third class
- Lacking in recall / F1

### Solution?

Try different models with different loss functions and see which one meets our design goals.



Source: <https://xkcd.com/1838/>

## The Different Models We Tried

1

ResNet50 with Recall Loss

What is Recall Loss?  
What is Focal Loss?

$$L(x)_{Recall} = -\log\left(\frac{w_i e^{z_i}}{\sum_{j=1}^n e^{z_j}}\right)$$

Where  $w_i$  is the false negative rate of the model on class  $i$  on the current batch.

2

Wide ResNet50 2 with Recall Loss

What can these models do better than the first two models?



3

ResNet50 with Focal Loss

What's the difference between Wide ResNet50 2 and ResNet50?

$$FL(p_t) = -(1 - p_t)^\gamma \log(p_t)$$

Where  $p_t = \begin{cases} p, & \text{if } y = 1 \\ 1 - p, & \text{otherwise} \end{cases}$

And  $p$  is the model's estimate for the class with label  $y$ , the class identifier.



## Contender Model Results

1

ResNet50 with Recall Loss

	Precision (%)
NIL	87.08
MOD	68.34
SEV	46.03
	Recall (%)
NIL	74.08
MOD	75.15
SEV	54.50
	F1 (%)
NIL	79.60
MOD	71.03
SEV	48.32

2

Wide ResNet50 2 with Recall Loss

	Precision (%)
NIL	87.30
MOD	65.73
SEV	53.49
	Recall (%)
NIL	72.09
MOD	77.46
SEV	57.72
	F1 (%)
NIL	78.04
MOD	70.47
SEV	54.85

3

ResNet50 with Focal Loss

	Precision (%)
NIL	88.01
MOD	58.86
SEV	62.98
	Recall (%)
NIL	68.64
MOD	89.03
SEV	21.13
	F1 (%)
NIL	76.55
MOD	70.18
SEV	30.20

## Contender Model Results

1

ResNet50 with Recall Loss

	Precision (%)
NIL	87.08
MOD	68.34
SEV	46.03
	Recall (%)
NIL	74.08
MOD	75.15
SEV	54.50
	F1 (%)
NIL	79.60
MOD	71.03
SEV	48.32

2

Wide ResNet50 2 with Recall Loss

	Precision (%)
NIL	87.30
MOD	65.73
SEV	53.49
	Recall (%)
NIL	72.09
MOD	77.46
SEV	57.72
	F1 (%)
NIL	78.04
MOD	70.47
SEV	54.85



3

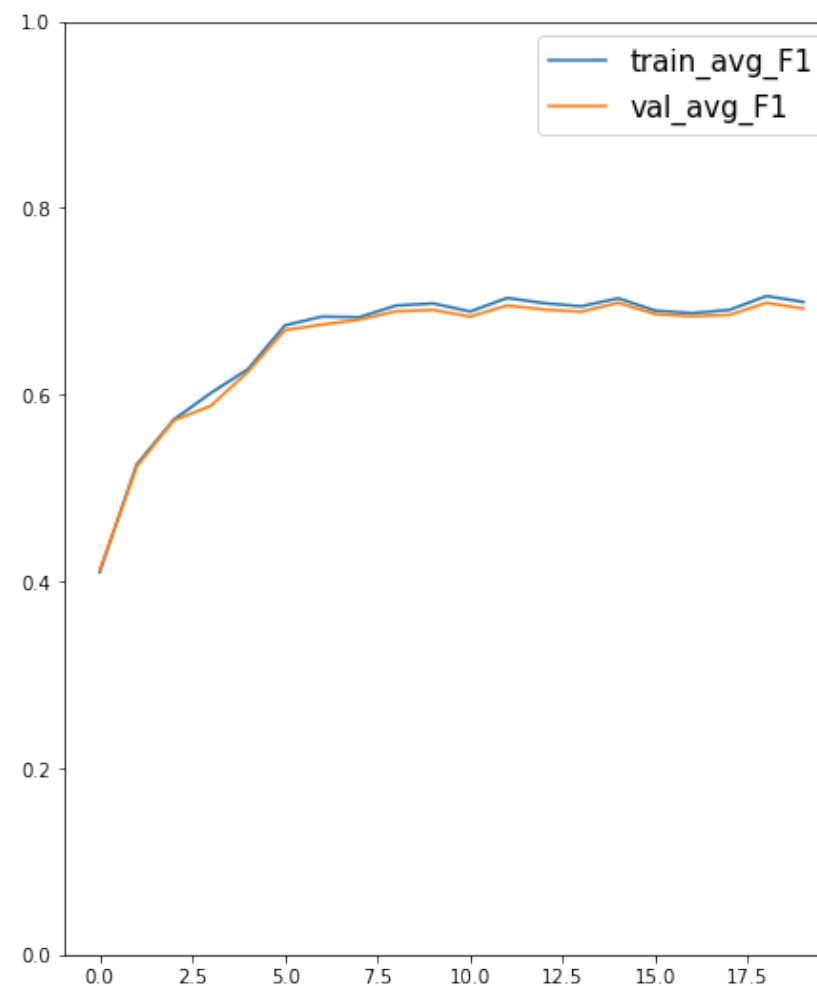

ResNet50 with Focal Loss

	Precision (%)
NIL	88.01
MOD	58.86
SEV	62.98
	Recall (%)
NIL	68.64
MOD	89.03
SEV	21.13
	F1 (%)
NIL	76.55
MOD	70.18
SEV	30.20

## Retrain the Chosen Model

- Wide ResNet50 2 with Recall Loss

	Precision (%)	Recall (%)	F1 (%)
NIL	88.53	74.52	80.28
Moderate	69.91	79.44	73.65
Severe	53.59	59.79	55.58



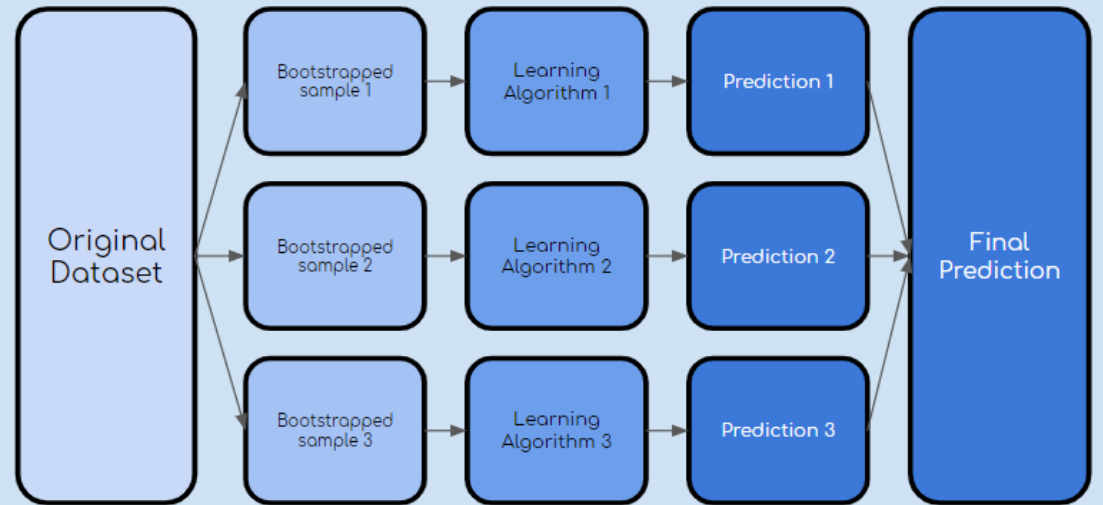
# Conclusion

---

## The Final Model – What Could've Been Done Differently

- Ensemble techniques
  - Bootstrap aggregating (Bagging)
- More Models
  - DenseNet, AlexNet, etc.
  - In general, CNNs are not rotation invariant
  - Customised model with max pooling?
- Domain-specific knowledge

### Ensemble Learning, Bagging, and Boosting



## Related Works and Information

- Recall Loss: <https://openreview.net/pdf?id=SlprFTIQP3>
- Focal Loss: <https://arxiv.org/pdf/1708.02002.pdf>



 Thank You For Listening! 

Any Questions?